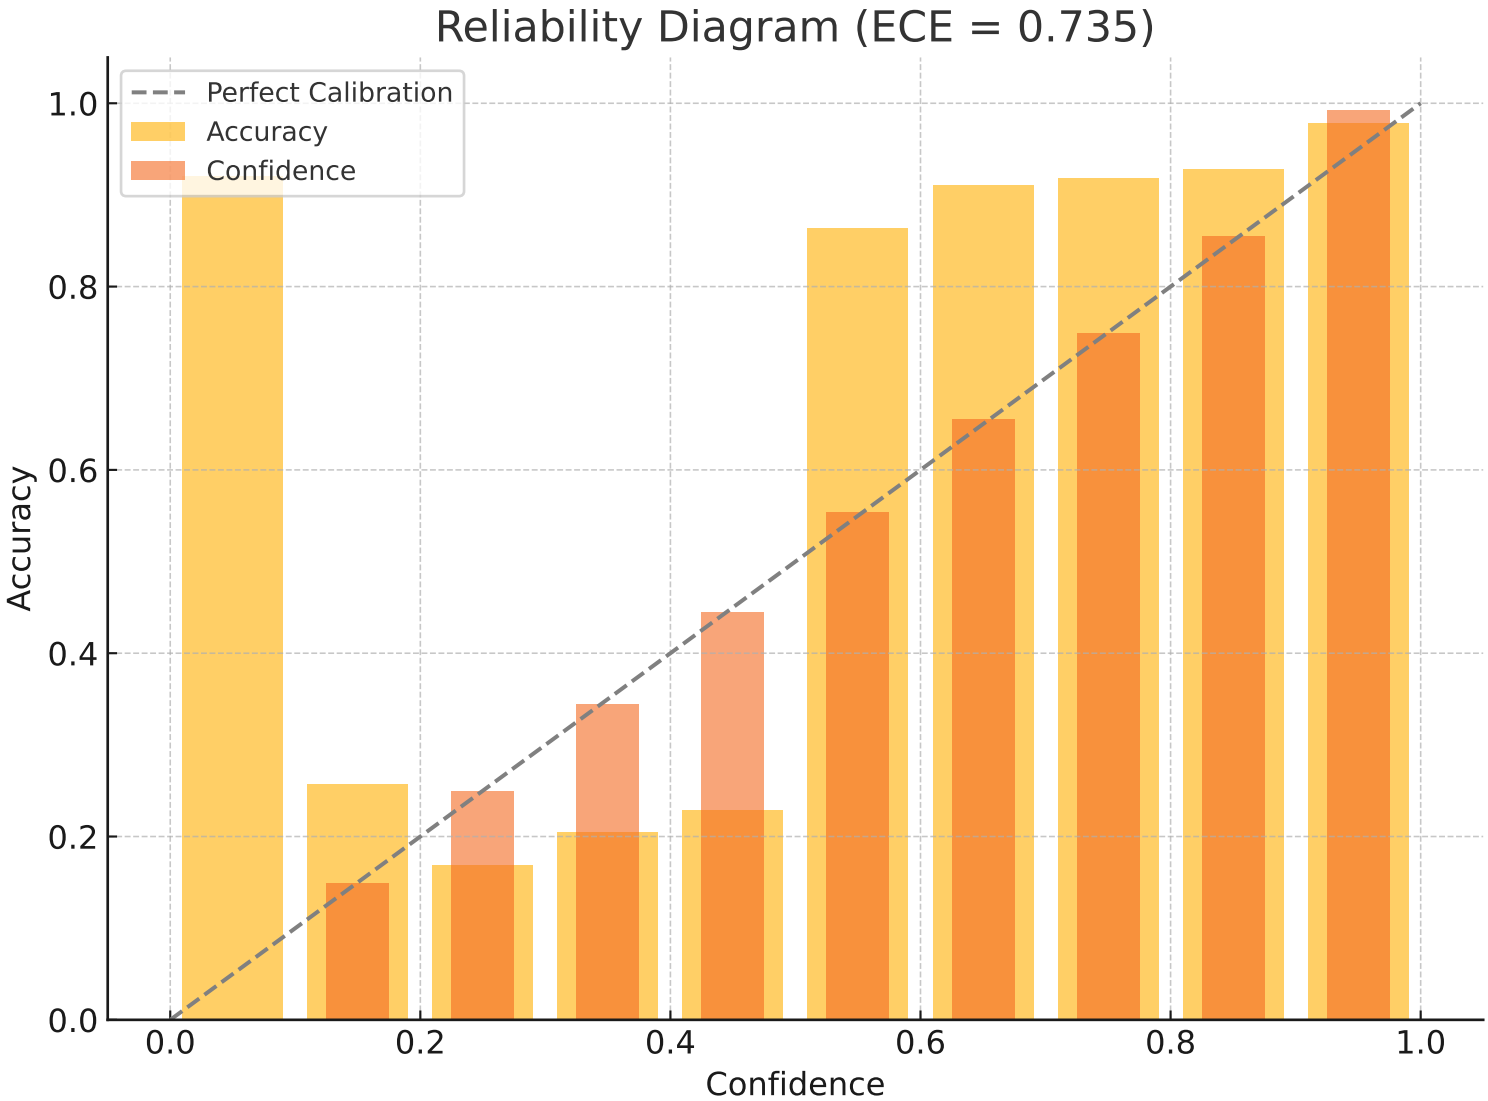
Reliability Diagram (ECE = 0.735)
Perfect Calibration
Accuracy
Confidence
Accuracy
Confidence
0.0
0.2
0.4
0.6
0.8
1.0

Uncertainty & Calibration Report
================================

Dataset Summary:

- Input: demo_cipo.csv
- Total samples: 26546
- Classes: Binary (Green / No)

Model Calibration Summary:

- Expected Calibration Error (ECE): 0.735
- Number of confidence bins: 10

Calibration Analysis:

The model shows signs of poor calibration (ECE = 0.735). Key observations:
1. Some low-confidence bins (e.g., 0.0–0.1) exhibit high accuracy, indicating underconfi
2. Some mid-confidence bins (e.g., 0.4–0.6) display mismatches between predicted confide
3. Overall, the model tends to be overconfident in several ranges, with a large divergen

Recommendations:

- Apply temperature scaling as a simple post-hoc calibration fix.
- Evaluate entropy-based or margin-based uncertainty metrics.
- Consider training a meta-model to predict model loss or uncertainty.
- Re-evaluate the effect of label smoothing or data augmentations.

Prepared by: Matan Hoffman